

**A LIGHTWEIGHT FRAMEWORK FOR
PRIVACY-PRESERVING
BEHAVIORAL AUTHENTICATION:
BALANCING RECOGNITION
ACCURACY AND SYSTEM LATENCY**

A PROJECT PROPOSAL SUBMITTED BY

W.M.T.R.S Weerakoon
(S20545)

to the

**DEPARTMENT OF STATISTICS AND COMPUTER
SCIENCE**

in partial fulfillment of the requirements of the

Degree of Bachelor of Science (Honours) in Computer Science
of the

UNIVERSITY OF PERADENIYA, SRI LANKA

2 February, 2026

Preface

This is a proposal for the Project in Computer Science II (Research Project – CSC4996) for partial fulfillment of the requirements of the Degree of Bachelor of Science (Honours) in Computer Science at the Department of Statistics and Computer Science, University of Peradeniya, Sri Lanka.

This proposal provides the scope and context of the research project to be undertaken. It details the aims and research questions, background, methodology, and project design. This document also provides a schedule for the completion of the project, including a description of anticipated results and final products

The intended audience of this document is the academic staff of the Department of Statistics and Computer Science, University of Peradeniya, who will evaluate the project to determine whether it should be approved as proposed, approved with modifications, or not approved.

Table of Contents

PREFACE	ii
TABLE OF CONTENTS	iii
LIST OF FIGURES	v
LIST OF TABLES	vi
LIST OF EQUATIONS	vii
LIST OF ABBREVIATIONS	viii
1.0 INTRODUCTION	1
1.1 Problem Statement	1
1.2 Research Aim and Objectives	2
1.3 Research Questions	2
2.0 BACKGROUND	4
2.1 Technical and Theoretical Background	4
2.1.1 Keystroke Dynamics	4
2.1.2 Mouse Dynamics	4
2.1.3 Recurrent Neural Networks (RNNs) & LSTMs	4
2.1.4 Johnson-Lindenstrauss (JL) Lemma	4
2.1.5 Homomorphic Encryption (HE)	5
2.2 Literature Review	5
2.2.1 Foundational Studies (Static Authentication)	5
2.2.2 Feature Engineering and Adaptive Systems	5
2.2.3 Mouse Dynamics and Multimodal Fusion	5
2.2.4 Deep Learning and Continuous Authentication	6
2.2.5 Scalability and System Performance	6
2.3 Research Gap	6
2.4 Assumptions and Constraints	7
3.0 METHODOLOGY AND PROJECT DESIGN	9
3.1 Overview of the Proposed Methodology/Research Design	9
3.2 Data Collection	10
3.3 Ethical Considerations	10
3.4 Evaluation and Validation	10
4.0 ANTICIPATED RESULTS/FINAL PRODUCTS	10
4.1 Expected Outcomes	10
4.2 Project Deliverables	11
4.3 Project Timeline	11
5.0 PROJECT SCHEDULE	11
5.1 Timeline and Gantt Chart	11
5.2 Milestones and Deliverables	11
REFERENCES	12

APPENDICES	13
-----------------------------	-----------

List of Figures

List of Tables

1	Proposed Project Timeline	11
---	-------------------------------------	----

List of Equations

List of Abbreviations

1.0 Introduction

Overview of the Research Domain In mechanisms such as Personal Identification Numbers (PINs) and passwords are increasingly becoming single points of failure [1]. They are susceptible to vulnerabilities like shoulder surfing, brute-force attacks, and social engineering [1], [2]. To address these vulnerabilities, Behavioral Biometrics—specifically Keystroke Dynamics and Mouse/Touch Dynamics—has emerged as a powerful alternative [3], [4]. Unlike static passwords, behavioral biometrics allow for continuous, passive authentication, verifying a user’s identity based on *how* they interact with a device rather than *what* they know [5], [6].

1.1 Problem Statement

The primary problem addressed by this research is the critical trade-off between **user privacy** and **system latency** in behavioral authentication systems on resource-constrained mobile devices. While behavioral biometrics—such as keystroke dynamics and touch interactions—offer a robust solution for continuous authentication, the storage and processing of these behavioral patterns present significant security risks [3].

This issue affects millions of mobile banking and enterprise users who rely on smartphones for sensitive transactions. Unlike passwords, behavioral biometrics are immutable; a user cannot change their typing rhythm or hand geometry if the biometric template is compromised. Therefore, a breach of raw behavioral data constitutes a permanent loss of digital identity [1].

Current approaches fail to address this problem effectively due to a technical dichotomy between accuracy and efficiency:

- **Privacy Gaps in High-Accuracy Models:** State-of-the-art deep learning frameworks, such as those utilizing Recurrent Neural Networks (RNNs) or Image-based Encoding (e.g., KDPrint), achieve low Equal Error Rates (EER) but typically require the storage of **raw behavioral features** [2], [4]. This creates a single point of failure where the database becomes a high-value target for attackers.
- **Efficiency Gaps in Privacy-Preserving Methods:** Conversely, strong cryptographic solutions like Fully Homomorphic Encryption (FHE) allow for secure computation but suffer from prohibitive computational overhead. Research indicates that such heavy encryption schemes often introduce latencies ranging from seconds to minutes, rendering them impractical for **real-time, continuous authentication** where decisions must be made in milliseconds to maintain a seamless user experience [3].

There is currently no lightweight framework that effectively balances these conflicting requirements. This research seeks to bridge this gap by utilizing **Orthogonal Random Projections** to secure templates without degrading the speed or accuracy required for mobile devices.

1.2 Research Aim and Objectives

The primary aim of this research is to develop a lightweight, privacy-preserving framework for continuous behavioral authentication on mobile devices. This framework seeks to balance recognition accuracy and system latency by utilizing Orthogonal Random Projections for template security and Deep Support Vector Data Description (Deep SVDD) for efficient anomaly detection.

1. **To design a privacy-preserving feature transformation pipeline:** Develop a mechanism using **Orthogonal Random Projections (Johnson-Lindenstrauss Lemma)** [7] that secures behavioral biometric templates (making them mathematically irreversible) while preserving the Euclidean distances required for accurate pattern recognition.
2. **To optimize feature engineering for mobile efficiency:** Implement **KDPrint-style standardization** to transform raw time-series data into standardized image encodings, ensuring high recognition accuracy without the noise sensitivity of Min-Max scaling [4].
3. **To implement a lightweight anomaly detection model:** Develop a **Deep SVDD (Support Vector Data Description)** classifier [8] capable of running offline on mid-range mobile devices to distinguish between genuine users and imposters with minimal computational overhead.
4. **To evaluate the trade-off between privacy, accuracy, and latency:** Conduct a comparative analysis of the proposed framework against existing baselines (such as raw-data RNNs and Homomorphic Encryption), measuring performance metrics including **Equal Error Rate (EER)**, **System Latency (ms)**, and **Memory Usage** [2], [8].

1.3 Research Questions

To address the identified gaps in privacy and efficiency, this research aims to answer the following key questions:

1. **Primary Research Question:** To what extent can **Orthogonal Random Projections** (based on the Johnson-Lindenstrauss Lemma) balance

the conflicting requirements of template privacy, recognition accuracy, and system latency in behavioral authentication?

2. **Impact on Privacy and Irreversibility:** How effective is the proposed projection mechanism in rendering behavioral templates mathematically irreversible to attackers, compared to storing raw features or using standard Min-Max scaling?
3. **Feasibility for Mobile Environments:** Can a **Deep SVDD** (Support Vector Data Description) anomaly detection model achieve real-time authentication latency (e.g., < 100ms) on mid-range mobile devices without exceeding memory constraints?

2.0 Background

2.1 Technical and Theoretical Background

2.1.1 Keystroke Dynamics

This is the measurement of biomechanical typing patterns. The fundamental features include Dwell Time (duration a key is pressed) and Flight Time (latency between releasing one key and pressing the next). These features form a unique "digital signature" for each user.

2.1.2 Mouse Dynamics

This involves analyzing the unique behavioral patterns of a user's mouse interactions. Unlike simple click-tracking, this research focuses on complex motor-skill features:

- **Movement Efficiency:** The ratio of the straight-line distance to the actual path taken (analyzing hand jitter and curvature).
- **Velocity & Acceleration Profiles:** The rate of speed change as the cursor approaches a target (e.g., users often decelerate differently when clicking a button).
- **Click-to-Click Latency:** The timing between releasing a button and moving to the next location.
- **Drag-and-Drop characteristics::** The pressure and speed consistency during sustained click events.

2.1.3 Recurrent Neural Networks (RNNs) & LSTMs

Since keystroke data is inherently sequential (a time-series of events), standard feed-forward networks often fail to capture temporal dependencies. RNNs, and specifically Long Short-Term Memory (LSTM) networks, are theoretically suited for this task as they maintain a "memory" of previous inputs, allowing them to model complex, non-linear typing rhythms over time.

2.1.4 Johnson-Lindenstrauss (JL) Lemma

To address the "curse of dimensionality" and system latency, this research utilizes the JL Lemma. This mathematical theorem states that points in a high-dimensional space can be projected into a lower-dimensional space using Orthogonal Random Projections while approximately preserving the Euclidean distances

between them. This allows for lightweight processing without significant loss of accuracy.

2.1.5 Homomorphic Encryption (HE)

To ensure privacy, the system employs Homomorphic Encryption, a cryptographic form that allows computations to be performed on encrypted data without first decrypting it. This ensures that the user's raw biometric template is never exposed in plaintext during the authentication process.

2.2 Literature Review

2.2.1 Foundational Studies (Static Authentication)

Early research laid the groundwork for typing pattern analysis. Gaines et al. (1980) demonstrated that typing rhythms are unique to individuals, identifying that even simple digraph latencies could distinguish users with high confidence. Building on this, Joyce and Gupta (1990) formalized the use of "latency signatures" for identity verification, showing that comparing a test signature against a mean reference signature could achieve low impostor pass rates. Monroe et al. (1999) extended this concept to "password hardening," combining typing rhythms with passwords to increase entropy against offline attacks.

2.2.2 Feature Engineering and Adaptive Systems

As research matured, the focus shifted to handling the variability in human behavior. Kim et al. (2018) highlighted the limitations of fixed-text authentication and proposed "user-adaptive feature extraction." Their work demonstrated that by dynamically selecting features based on a user's specific typing speed ranks (rather than a fixed keyboard layout), the Equal Error Rate (EER) could be significantly reduced. This underscored the need for personalized models in behavioral authentication.

2.2.3 Mouse Dynamics and Multimodal Fusion

Research into mouse dynamics has evolved from simple statistics to complex trajectory analysis. Ahmed and Traore (2007) pioneered the use of movement speed and direction histograms, achieving reasonable accuracy but noting high false-positive rates in unconstrained environments. Later, Zheng et al. (2011) improved this by analyzing "point-by-point" angle-based metrics, proving that fine-grained motor skills are harder to forge than simple click statistics.

Recent studies have shifted toward Multimodal Authentication, combining keystroke and mouse data to overcome the weaknesses of single-modality systems. Mondal and Bours (2017) demonstrated that fusing these two biometrics significantly reduces the Equal Error Rate (EER) because an imposter is unlikely to mimic both a user’s typing rhythm and their mouse hand-eye coordination simultaneously. However, existing multimodal frameworks often simply concatenate feature vectors, leading to massive dimensionality that slows down real-time processing—a problem this research aims to solve using the JL Lemma.

2.2.4 Deep Learning and Continuous Authentication

Recent approaches have adopted Machine Learning (ML) and Deep Learning (DL) to improve accuracy. Zareen et al. (2018) utilized Bayesian Regularized Feed-Forward Neural Networks, achieving an EER of 0.9%, which outperformed many standard statistical methods. Kiyani et al. (2020) proposed a ”Robust Recurrent Confidence Model” (R-RCM) using ensemble learning. Their system verified identity on a per-action basis (Continuous Authentication), validating the effectiveness of sequential models for real-time monitoring.

2.2.5 Scalability and System Performance

While accuracy has improved, scalability remains a challenge. Islam et al. (2021) introduced the notion of ”Scalable Behavioral Authentication,” analyzing how verification error rates increase as the user database grows. They proposed a ”doppelganger-based” personalized verification algorithm to mitigate error growth, highlighting that system performance must be evaluated not just on accuracy, but on its ability to handle large-scale deployments.

2.3 Research Gap

Despite the extensive literature on improving the accuracy of keystroke dynamics (Kiyani et al., 2020; Zareen et al., 2018), a critical trilemma remains unsolved: balancing Accuracy, Efficiency, and Privacy.

1. **Computation vs. Latency in Multimodal Systems:** Integrating Mouse Dynamics with Keystroke Dynamics doubles the feature space complexity. Processing this high-dimensional, multimodal stream (typing + mouse trajectories) using deep models (RNNs) creates a bottleneck that current systems fail to address efficiently.
2. **Privacy Vulnerability:** Most existing frameworks (Kim et al., 2018; Joyce, 1990) focus on verifying raw feature vectors. They lack robust ”template

protection” mechanisms. If the central database is compromised, the user’s raw behavioral biometrics are lost forever.

3. **Lack of Integrated Privacy-Preserving Architectures:** While encryption exists, standard encryption prevents the system from performing the distance calculations needed for authentication (like Euclidean distance or KL-divergence mentioned in Islam et al.).

Gap Definition: There is currently no unified framework that utilizes Orthogonal Random Projections (JL Lemma) to compress the combined feature space of both Keystroke and Mouse Dynamics for lightweight processing, while simultaneously preserving privacy using Homomorphic Encryption. This study aims to bridge this gap by creating a fast, privacy-preserving, multimodal authentication system.

2.4 Assumptions and Constraints

- **Assumptions:**
 - It is assumed that the user’s typing behavior is relatively stable over short periods but may exhibit gradual ”concept drift” which the model must accommodate.
 - It is assumed that the users are utilizing standard physical keyboards; virtual/touchscreen keyboards are outside the scope of this specific study (unless specified otherwise).
 - It is assumed that the mouse data collection frequency (e.g., 50Hz) is sufficient to capture micro-movements without overwhelming the system bus.
 - The ”trust” of the endpoint device (personal laptop) is assumed for the initial data capture phase before encryption.
- **Constraints:**
 - **Hardware Limitations:** The proposed model must be lightweight enough to run on standard consumer hardware (e.g., a laptop with a mid-range GPU like an RTX 3050) without causing noticeable input lag.
 - **Data Availability:** The research is constrained by the need to collect a custom dataset or use public datasets that may not perfectly reflect the specific ”free-text” behavior required for continuous authentication.

- **Encryption Overhead:** Homomorphic Encryption introduces significant computational overhead. The system is constrained to optimize this trade-off to ensure the authentication decision happens within a usable timeframe (e.g., under 200ms per window).

3.0 Methodology and Project Design

Rationale for Chosen Methods:

- **Multimodal Fusion:** Single-modality systems (keystroke only) are prone to mimicry attacks. Fusion with mouse dynamics increases the entropy of the user profile, making forgery exponentially harder.
- **Deep SVDD & LSTMs:** Unlike static classifiers (e.g., SVM), Long Short-Term Memory (LSTM) networks are selected for their ability to model the temporal dependencies in sequential data. Deep Support Vector Data Description (Deep SVDD) is chosen as the anomaly detector because it is a "one-class" classifier, meaning it can train on only the legitimate user's data without requiring impostor data during the training phase.
- **Johnson-Lindenstrauss (JL) Lemma:** To counter the computational overhead of Homomorphic Encryption, the JL Lemma is applied to project high-dimensional biometric feature vectors into a lower-dimensional space. This allows for faster encrypted computations with mathematically guaranteed distance preservation.

3.1 Overview of the Proposed Methodology/Research Design

The proposed system architecture is designed as a post-acquisition processing pipeline. It assumes the availability of raw timestamped log files containing keystroke and mouse events. The methodology follows a sequential four-phase workflow:

1. **Feature Extraction & Temporal Fusion:** The system ingests raw event logs and converts them into synchronized time-series feature vectors.
 - **Keystroke Features:** Extraction of Flight Time (latency between KeyUP $n \rightarrow$ KeyDOWN $n + 1$) and *Dwell Time* (duration of KeyDOWN $n \rightarrow$ KeyUP n).
 - **Mouse Features:** Calculation of higher-order motor metrics including Velocity Profiles, Angular Velocity, and Curvature Distance Ratio (efficiency of movement).
 - **Multimodal Fusion:** The two independent streams are aligned using sliding time windows (e.g., $t = 10s$) to create unified "behavioral frames" representing the user's complete interaction state.

2. **Dimensionality Reduction (JL Layer):** To mitigate the "curse of dimensionality" caused by fusing two biometric streams, the high-dimensional fused vector (d) is projected onto a lower-dimensional subspace (k , where $k \ll d$) using the Johnson-Lindenstrauss (JL) Lemma. This is achieved by multiplying the feature vector by a sparse random matrix (R) to produce a compressed, privacy-hardened embedding.
3. **Privacy-Preserving Transformation:** The compressed embeddings are encrypted using a Partially Homomorphic Encryption (PHE) scheme (e.g., Paillier). This ensures that all subsequent distance calculations required for authentication are performed in the encrypted domain, preventing the exposure of the user's behavioral template.
4. **Anomaly Detection (Deep SVDD):** The encrypted vectors are fed into a Deep Support Vector Data Description (Deep SVDD) model. This model learns a compact hypersphere boundary encapsulating the legitimate user's "normal" behavior. During the verification phase, any input vector falling outside this learned boundary is flagged as an anomaly (potential impostor) without ever decrypting the data.

3.2 Data Collection

3.3 Ethical Considerations

3.4 Evaluation and Validation

- **False Acceptance Rate (FAR)**
- **False Rejection Rate (FRR)**
- **Equal Error Rate (EER)**
- **System Latency (ms)**

4.0 Anticipated Results/Final Products

4.1 Expected Outcomes

- A fully functional Android application capable of continuous authentication.
- A comparative analysis report on different RNN architectures (LSTM vs. GRU).
- A research paper submitted to an IEEE conference.

4.2 Project Deliverables

1. **Software Prototype:** Source code and executable APK.
2. **Final Thesis:** Comprehensive documentation of the research process.
3. **User Manual:** Guide for installing and testing the application.

4.3 Project Timeline

5.0 Project Schedule

5.1 Timeline and Gantt Chart

Phase	Activity	Duration
1	Literature Review	Months 1-2
2	Data Collection	Month 3
3	Model Implementation	Months 4-5
4	Testing and Validation	Month 6
5	Thesis Writing	Months 7-8

Table 1: Proposed Project Timeline

5.2 Milestones and Deliverables

- **Milestone 1:** Completion of Literature Review (Date)
- **Milestone 2:** Prototype Development (Date)
- **Final Submission:** Thesis and Defense (Date)

References

- [1] S. S. Pirzado and et al., “Keystroke dynamics based technique to enhance the security in smart devices,” *Mehran University Research Journal of Engineering and Technology*, vol. 40, no. 1, pp. 123–133, 2021.
- [2] A. T. Kiyani, A. Lasebae, K. Ali, M. U. Rehman, and B. Haq, “Continuous user authentication featuring keystroke dynamics based on robust recurrent confidence model and ensemble learning approach,” *IEEE Access*, vol. 8, pp. 156 177–156 189, 2020.
- [3] M. Rahman and et al., “Scalable behavioral authentication for mobile devices,” *Sensors*, vol. 21, no. 14, p. 4890, 2021.
- [4] Y. Kim, D. Shin, and N. Kwon, “Kdprint: Passive authentication using keystroke dynamics-to-image encoding via standardization,” *arXiv preprint arXiv:2405.01080*, 2024, Accessed via arXiv.
- [5] S. J. Shepherd, “Continuous authentication by analysis of keyboard typing characteristics,” in *European Convention on Security and Detection*, IET, 1995, pp. 111–114.
- [6] F. Monrose and A. D. Rubin, “Password hardening based on keystroke dynamics,” in *Proceedings of the 6th ACM Conference on Computer and Communications Security (CCS)*, ACM, 1999, pp. 73–82.
- [7] W. B. Johnson and J. Lindenstrauss, “Extensions of lipschitz mappings into a hilbert space,” in *Conference in modern analysis and probability*, Contemporary Mathematics, American Mathematical Society, 1984, pp. 189–206.
- [8] L. Ruff et al., “Deep one-class classification,” in *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018, pp. 4393–4402.

Appendices

Appendix A: Survey Form

Appendix B: Code Snippets